

Hardware Project #5

Pipelined CPU

due at 5pm, Tue 5 Dec 2017

1 Purpose

In this project, you will be adapting your Hardware Project 4 code to build a pipelined processor. Each pipeline register will be represented by its own struct; for instance, the ID/EX pipeline register is represented by the struct `ID_EX`. Just like in a real processor, each pipeline register must contain all of the information that is saved from one clock cycle to another; there will **not** be any global Control struct (as there was in Hardware Project 4).

The core of each pipeline phase is an `execute_*` function, which typically will read from the pipeline register on the left, and write to the pipeline register on the right. However, the details vary from phase to phase - so read the details below.

1.1 Required Filenames to Turn in

Name your C file `proj_hw05.c`.

2 Required Instructions

Every student must implement the following instructions:

- `add`, `addu`, `sub`, `subu`, `addi`, `addiu`
(Treat the 'u' instructions exactly the same as their more ordinary counterparts. Just use C addition and subtraction, using signed integers. Ignore overflow.)
- `and`, `or`, **`nor`**
- `slt`, `slti`
- `lw`, `sw`
- `beq`, `bne`, `j`
- **`andi`**, **`ori`**
- **`lui`**

(The instructions in red above will require you to expand the standard CPU design by adding extra control bits, or extra values to existing controls.)

2.1 Extended Control Values

Just like in Hardware Project 4, we'll need to modify the standard CPU design a bit. You **must** add the following feature:

- Set `ALUsrc=2` to indicate **zero-extended imm16**

(This will be enough to pass the testcases that print out lots of debug data.)

In addition, as noted above, you must support other instructions as well - but for those, you can decide how to implement them. While you can expect that I will test them, I won't print out debug data when I do.

3 Data Forwarding, 1w Hazards, and Stalls

Your CPU must implement data forwarding. This will be handled in the EX phase, in the `EX_ALUinput*` functions. This means that, in almost all scenarios, it will not be necessary to stall the pipeline.

However, your processor must implement **1w** hazard detection in the ID phase; if you detect a hazard between the current instruction (in ID) and a **1w** instruction in front of it (in EX), then you will signal that a stall is required.

If you ask for a stall, my testcase code will handle what happens in IF; I will ensure that, on the next clock cycle, the ID phase will see the exact same instruction a second time. However, you will need to implement the stall in the ID/EX pipeline register; to do this, you must set all of the bits in that register to 0.

4 Branches

Your CPU will implement conditional and unconditional branches. These will be resolved in the ID phase; you will indicate when/if a branch is required, as well as the destination of the branch (see the details below).

Of course, in a real CPU, there would already be an instruction in the IF phase when a branch occurs. Two strategies are classically used; either the instruction in IF is flushed (causing a NOP in the instruction stream), or else a **branch delay slot** is used.

In our CPU simulation, we'll use a third, simpler method - it isn't technically correct, but it's easy to understand: we will "magically" fetch the next instruction (at the branch destination), and have it ready for execution in the very next clock cycle.

However, you will be required to handle the ID/EX pipeline register. Since you will be handling all of the branch logic in ID, this means that every branch instruction (conditional or unconditional) should be converted to a NOP when it reaches the ID/EX pipeline register¹.

¹If you choose to go beyond the required project, and implement the **JAL** instruction, then you will need to do some sort of operation for **JAL**. But this is not necessary for any of the required instructions.

5 The Phases

Generally, each phase reads from the pipeline register on the left, and writes to the pipeline register on the right. The testcase will typically have two copies of each pipeline register: the “old” and “new” values. The old values are the contents of the register at the beginning of this clock cycle, and the new values are the contents of the register at the end. Thus, you may write to the “new” copy of each register, without worrying that you might be overwriting critical data in the “old,” which might be needed by another pipeline phase during this cycle.

5.1 IF

The IF phase will be implemented by my testcases - however, your code in ID will instruct it what to do.

5.2 ID

The ID phase is the most complex of all, because it has to implement both stall and branch/jump logic. For this reason, there are several different function calls which will occur:

- `extract_instructionFields()`

This works exactly like Hardware Project 4.

- `IDtoIF_get_stall()`

This function asks if a stall is required. If you need to stall the ID phase, return 1; if not, then return 0.

The parameters are the Fields of the current instruction, plus the ID/EX pipeline register, for the instruction ahead. This function **must not modify any of these fields** - simply query to determine if a stall is required.

If a stall is required, then the IF phase will also stall, and so you will see the exact same instruction on the next clock cycle.

This function must pay attention to the opcode and funct, since different instructions use different registers; only stall when the `lw` actually impacts one of the inputs to this instruction.

If you don't recognize the opcode or funct (meaning that this is an invalid instruction), return 0 from this function.

- `IDtoIF_get_branchControl()`

This asks the ID phase if the current instruction (in ID) needs to perform a branch/jump. The parameters are the Fields for this instruction, along with the `rsVal` and `rtVal` for this instruction.

If you return 0, then the PC will advance as normal. (If you ask for a stall, then you must also return this `branchControl` value.)

If you return 1, then the PC will jump to the (relative) branch destination - see `calc_branchAddr()`.

If you return 2, then the PC will jump to the (absolute) jump destination - see `calc_jumpAddr()`.

If you return 3, then the PC will jump to `rsVal`. (You will not need to use this feature unless you decide to add support for `JR`.)

If you don't recognize the opcode or funct (meaning that this is an invalid instruction), return 0 from this function.

- `calc_branchAddr()`

This asks you to calculate the address that you would jump to if you perform a conditional branch (`BEQ`, `BNE`). This function should model a simple branch adder in hardware - and thus, it will calculate this value on **every clock cycle, and for every instruction** - even if there is no possible way that it might be used.

Essentially, this is one of 4 inputs to a MUX - where the MUX is controlled by the `branchControl` value above. (My code in the IF phase will implement this MUX.)

- `calc_jumpAddr()`

This asks you to calculate the address that you would jump to if you perform an unconditional branch (`J`²).

As with `calc_branchAddr()`, this must calculate this value on every clock cycle, and for every instruction - even if there is no possible way that it might be used.

- `execute_ID()`

This function implements the core of the ID phase. Its first parameter is the `stall` setting (exactly what you returned from `IDtoIF_get_stall()`). The next is the Fields for this instruction, followed by the `rsVal` and `rtVal`; last is a pointer to the (new) ID/EX pipeline register.

Decode the opcode and funct, and set all of the fields of the `ID_EX` struct. (I don't define how you might use the `extra*` fields.)

As in Hardware Project 4, you will return 1 if you recognize the opcode/funct; return 0 if it is an invalid instruction.

5.3 EX

The EX phase includes the following functions:

- `EX_getALUinput1()`

²This would also be used for the `JAL` instruction, if you decided to add support for that.

This function must return the value which should be delivered to input 1 of the ALU. The first parameter is the current ID/EX register; it also has pointers to the current EX/MEM and MEM/WB registers.

This function must not modify any of the structs; it just returns a 32-bit value.

- `EX_getALUinput2()`

This is the same function, but for ALU input 2. It must also handle immediate values (see the `ALUsrc` control value).

- `execute_EX()`

This function implements the core of the EX phase. It has a pointer to the ID/EX pipeline register (which it must not modify), the two input values (see above), and a pointer to the EX/MEM pipeline register (which it must fill).

This phase must choose between the two possible destination registers (`rt,rd`); store the chosen register into `writeReg` (a 5-bit field). Notice that the `regWrite` control (1 bit) indicates **whether** we will write to a register, and `writeReg` indicates **which** we will write to.

5.4 MEM

The MEM phase only includes a single function:

- `execute_MEM()`

This function works more or less like `execute_MEM()` from Hardware Project 4; it may read or write memory.

5.5 WB

The WB phase only includes a single function:

- `execute_WB()`

This function works more or less like `execute_WB()` from Hardware Project 4; it may update a register.

6 A Note About Grading

Your code will be tested automatically. Therefore, your code must:

- Use exactly the filenames that we specify (remember that names are case sensitive).
- **Not** use any other files (unless allowed by the project spec) - since our grading script won't know to use them.

- Follow the spec precisely (don't change any names, or edit the files I give you, unless the spec says to do so).
- (In projects that require output) match the required output **exactly!** Any extra spaces, blank lines misspelled words, etc. will cause the testcase to fail.

To make it easy to check, I have provided the grading script. I **strongly recommend** that you download the grading script and all of the testcases, and use them to test your code from the beginning. You want to detect any problems early on!

6.1 Testcases

You can find the project files for this project (including any testcases I've provided) at:

http://lecturer-russ.appspot.com/classes/cs252/fall17/projects/proj_hw05/

You can also find them on Lectura at

/home/russell11/cs252.website/projects/proj_hw05/

For assembly language programs, the testcases will be named `test*.s`. For C programs, the testcases will be named `test*.c`. For Java programs, the testcases will be named `Test*.java`. (You will only have testcases for the languages that you have to actually write for each project, of course.)

Each testcase has a matching output file, which ends in `.out`; our grading script needs to have both files available in order to test your code.

For many projects, we will have "secret testcases," which are additional testcases that we do not publish until after the solutions have been posted. These may cover corner cases not covered by the basic testcase, or may simply provide additional testing. **You are encouraged to write testcases of your own, in order to better test your code.**

6.2 Automatic Testing

We have provided a testing script (in the same directory), named `grade_proj_hw04`. Place this script, all of the testcase files (including their `.out` files if assembly language), and your program files in the same directory. (I recommend that you do this on Lectura, or a similar department machine. It **might** also work on your Mac, but no promises!)

6.3 Writing Your Own Testcases

The grading script will grade your code based on the testcases it finds in the current directory. Start with the testcases I provide - however, I encourage you to write your own as well. If you write your own, simply name your testcases using the same pattern as mine, and the grading script will pick them up.

While you normally cannot share code with friends and classmates, **test-cases are the exception**. We encourage you to share your testcases - ideally by posting them on Piazza. Sometimes, I may even pick your testcase up to be part of the official set, when I do the grading!

7 Turning in Your Solution

You must turn in your code using D2L, using the Assignment folder for this project. Turn in only your program; do not turn in any testcases or other files.